

SQL

Data Definition Language

The SQL **data-definition language (DDL)** allows the specification of information about relations, including:

- ❑ The schema for each relation.
- ❑ The domain of values associated with each attribute.
- ❑ Integrity constraints
- ❑ And as we will see later, also other information such as
 - ❑ The set of indices to be maintained for each relations.
 - ❑ Security and authorization information for each relation.
 - ❑ The physical storage structure of each relation on disk.



Create Table Construct

? An SQL relation is defined using the **create table** command:

```
create table  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

? r is the name of the relation

? each A_i is an attribute name in the schema of relation r

? D_i is the data type of values in the domain of attribute A_i

? Example:

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null,  
    dept_name varchar(20),  
    salary     numeric(8,2))
```

? **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);

? **insert into** *instructor* **values** ('10211', null, 'Biology', 66000);



Integrity Constraints in Create Table

- ❑ **not null**
- ❑ **primary key** (A_1, \dots, A_n)
- ❑ **foreign key** (A_m, \dots, A_n) **references** r

Example: Declare ID as the primary key for *instructor*

```
.  
    create table instructor (  
        ID          char(5),  
        name        varchar(20) not null,  
        dept_name    varchar(20),  
        salary       numeric(8,2),  
        primary key (ID),  
        foreign key (dept_name) references department)
```

primary key declaration on an attribute automatically ensures **not null**



Drop and Alter Table Constructs

? **drop table** *student*

? Deletes the table and its contents

? **delete from** *student*

? Deletes all contents of table, but retains table

? **alter table**

? **alter table** *r* **add** *A D*

- ▶ where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
- ▶ All tuples in the relation are assigned *null* as the value for the new attribute.

? **alter table** *r* **drop** *A*

- ▶ where *A* is the name of an attribute of relation *r*
- ▶ Dropping of attributes not supported by many databases



Basic Query Structure

- ❑ The SQL **data-manipulation language (DML)** provides the ability to query information, and insert, delete and update tuples
- ❑ A typical SQL query has the form:

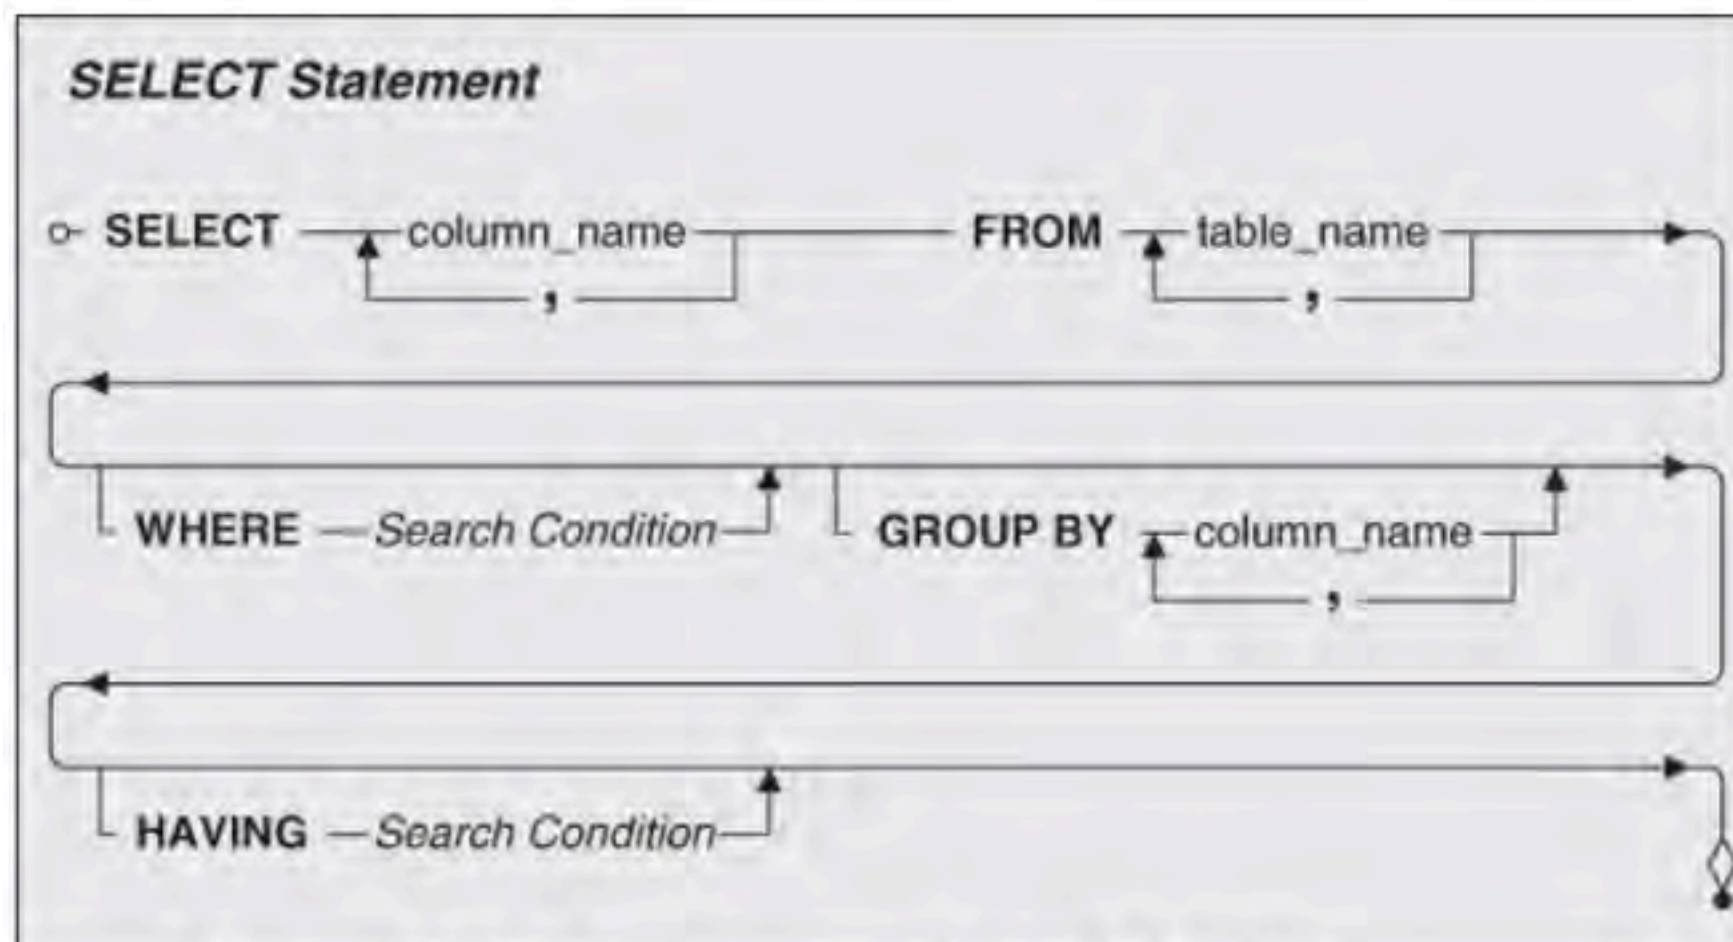
select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

- ❑ A_i represents an attribute
 - ❑ R_i represents a relation
 - ❑ P is a predicate.
- ❑ The result of an SQL query is a relation.



Select

- ❓ Forms the basis the basis of every question we pose to the DB
- ❓ Select –
 - ❓ Querying the DB.
 - ❓ Composed of several distinct keywords known as clauses
 - ▶ Some clauses are required, while others are optional
 - ▶ Each clause has one or more keywords that represent required or optional values



Select Clauses

- ❓ Select – Primary clause of the Select Statement (Required)
 - ❓ Specify columns you want in the result set of your query
 - ❓ Columns come from table or view
- ❓ FROM – Specify table or view (Required)
- ❓ Where – Used for filtering information returned (Optional)
- ❓ Group By -- Used to divide the information in distinct groups (optional)
 - ❓ When you use aggregate functions in the SELECT clause to produce summary information, you use the Group BY clause
- ❓ HAVING-- Filters the result of aggregate functions in grouped information (optional)
 - ❓ Similar to WHERE clause – HAVING clause is followed by an expression that evaluates to true, false, or unknown.



Select Cont.

- ❓ Information requested from the DB, is in the form of a question e.g.,
 - ❓ *“Which cities do our customers live in”*
 - ❓ *“Show me a current list of our employees and their phone numbers.”*
 - ❓ *What kind of classes do we currently offer?”*
 - ❓ *“Give me the names of the folks on our staff and the dates they were hired.”*
- ❓ You can translate the question into a formal request using the form:
 - ❓ Select <item> from the <source>
 - ❓ Replace words such as “Which, Show”, to SELECT
 - ❓ Identify nouns
 - ▶ Determine if noun represents the an item you want to see or
 - ▶ A Table that contains in which items are stored



Select Cont.

- *Which cities do our customers live in”*
 - *Which → SELECT*
 - *Cities → Items*
 - *Customer → Table*
- *Select City from the Customer table*
- Once you cleanup.. It looks like:
 - *Select city from the customers table*

SELECT Statement



Remove duplicates

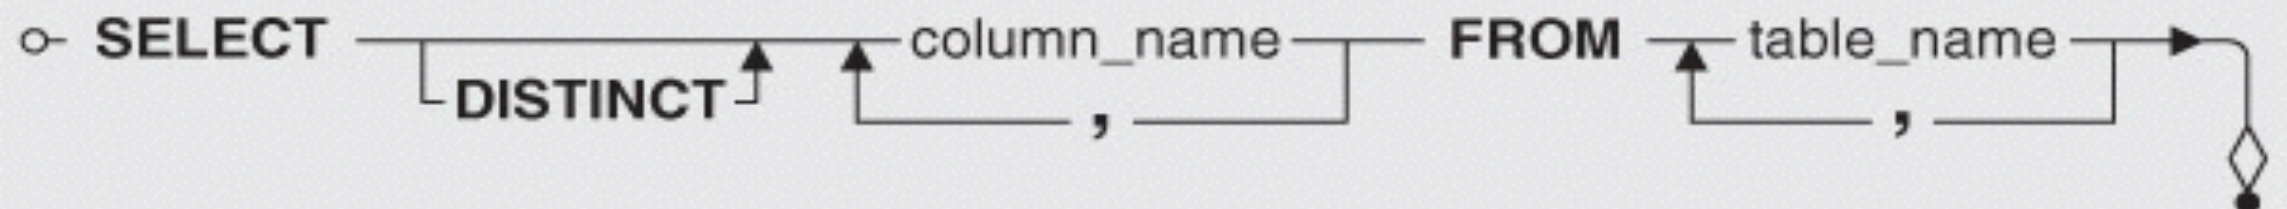
- ❓ SQL allows duplicates in relations as well as in query results.
- ❓ To force the elimination of duplicates, insert the keyword **distinct** after select.
- ❓ Find the names of all departments with instructor, and remove duplicates

```
select distinct dept_name  
from instructor
```

- ❓ The keyword **all** specifies that duplicates not be removed.


```
select all dept_name  
from instructor
```

SELECT Statement



Sorting Information

- ❓ Rows of the result set returned by a SELECT are unordered
- ❓ Result set is sorted by using ORDER BY clause



```
SELECT Category
FROM Classes
ORDER BY Category
```

```
SELECT VendName, VendZipCode
FROM Vendors
ORDER BY VendZipCode
```

```
SELECT VendName, VendZipCode
FROM Vendors
ORDER BY VendZipCode DESC
```



Joined Relations

- ❓ **Join operations** take two relations and return as a result another relation.
- ❓ A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- ❓ The join operations are typically used as subquery expressions in the **from** clause



Join operations – Example

? Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

? Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

? Observe that

prereq information is missing for CS-315 and
course information is missing for CS-437



Outer Join

- ❑ An extension of the join operation that avoids loss of information.
- ❑ Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- ❑ Uses *null* values.



Left Outer Join

? *course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

? *course*

? *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3



Right Outer Join

 *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



Full Outer Join

 *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



Joined Relations – Examples

[?] course inner join prereq on
course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

[?] What is the difference between the above, and a natural join?

[?] course left outer join prereq on
course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>



Joined Relations – Examples

[?] course natural right outer join prereq

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

[?] course full outer join prereq using (course_id)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



Nested Subqueries

- ❑ SQL provides a mechanism for the nesting of subqueries.
- ❑ A **subquery** is a **select-from-where** expression that is nested within another query.
- ❑ A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.



Example Query

- ❓ Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101

```
select count (distinct ID)  
from takes  
where (course_id, sec_id, semester, year) in  
        (select course_id, sec_id, semester, year  
         from teaches  
         where teaches.ID= 10101);
```

- ❓ **Note:** Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.



Modification of the Database

- ❑ Deletion of tuples from a given relation
- ❑ Insertion of new tuples into a given relation
- ❑ Updating values in some tuples in a given relation



Modification of the Database – Deletion

- ❑ Delete all instructors

delete from *instructor*

- ❑ Delete all instructors from the Finance department

delete from *instructor*

where *dept_name*= 'Finance';

- ❑ Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

delete from *instructor*

where *dept_name* in (**select** *dept_name*
from *department*
where *building* = 'Watson');



Deletion (Cont.)

- ❓ Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary) from instructor);
```

- ❓ Problem: as we delete tuples from deposit, the average salary changes
- ❓ Solution used in SQL:
 1. First, compute **avg** salary and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



Modification of the Database – Insertion

- ❓ Add a new tuple to *course*

insert into *course*

values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- ❓ or equivalently

insert into *course* (*course_id*, *title*, *dept_name*, *credits*)

values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- ❓ Add a new tuple to *student* with *tot_creds* set to null

insert into *student*

values ('3003', 'Green', 'Finance', *null*);



Insertion (Cont.)

- ❓ Add all instructors to the *student* relation with *tot_creds* set to 0

```
insert into student
  select ID, name, dept_name, 0
  from instructor
```

- ❓ The **select from where** statement is evaluated fully before any of its results are inserted into the relation (otherwise queries like
 insert into table1 select * from table1
 would cause problems, if *table1* did not have any primary key defined.



Modification of the Database – Updates

- ❑ Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise

- ❑ Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

- ❑ The order is important
- ❑ Can be done better using the **case** statement (next slide)



Built-in Data Types in SQL

[?] **date:** Dates, containing a (4 digit) year, month and date

[?] Example: **date** '2005-7-27'

[?] **time:** Time of day, in hours, minutes and seconds.

[?] Example: **time** '09:00:30' **time** '09:00:30.75'

[?] **timestamp:** date plus time of day

[?] Example: **timestamp** '2005-7-27 09:00:30.75'

[?] **interval:** period of time

[?] Example: **interval** '1' day

[?] Subtracting a date/time/timestamp value from another gives an interval value


[?] Interval values can be added to date/time/timestamp values



User-Defined Types

 **create type** construct in SQL creates user-defined type

create type *Dollars* **as numeric (12,2) final**

 **create table** *department*
(*dept_name* **varchar** (20),
building **varchar** (15),
budget *Dollars*);



Large-Object Types

- ❓ Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
- ❓ **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
- ❓ **clob**: character large object -- object is a large collection of character data
- ❓ When a query returns a large object, a pointer is returned rather than the large object itself.



Authorization

Forms of authorization on parts of the database:

- ❑ **Read** - allows reading, but not modification of data.
- ❑ **Insert** - allows insertion of new data, but not modification of existing data.
- ❑ **Update** - allows modification, but not deletion of data.
- ❑ **Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- ❑ **Index** - allows creation and deletion of indices.
- ❑ **Resources** - allows creation of new relations.
- ❑ **Alteration** - allows addition or deletion of attributes in a relation.
- ❑ **Drop** - allows deletion of relations.



Authorization Specification in SQL

- ❑ The **grant** statement is used to confer authorization
 - grant** <privilege list>
 - on** <relation name or view name> **to** <user list>
- ❑ <user list> is:
 - ❑ a user-id
 - ❑ **public**, which allows all valid users the privilege granted
 - ❑ A role (more on this later)
- ❑ Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- ❑ The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).



Privileges in SQL

❓ **select**: allows read access to relation, or the ability to query using the view

❓ Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *instructor* relation:

grant select on *instructor* to U_1 , U_2 , U_3

❓ **insert**: the ability to insert tuples

❓ **update**: the ability to update using the SQL update statement

❓ **delete**: the ability to delete tuples.

❓ **all privileges**: used as a short form for all the allowable privileges



Revoking Authorization in SQL

❓ The **revoke** statement is used to revoke authorization.

revoke <privilege list>

on <relation name or view name> **from** <user list>

❓ Example:

revoke select on *branch* **from** U_1, U_2, U_3

❓ <privilege-list> may be **all** to revoke all privileges the revokee may hold.

❓ If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.

❓ If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.

❓ All privileges that depend on the privilege being revoked are also revoked.



Roles

- ❑ **create role** instructor;
- ❑ **grant** *instructor* **to Amit**;
- ❑ Privileges can be granted to roles:
 - ❑ **grant select on** *takes* **to instructor**;
- ❑ Roles can be granted to users, as well as to other roles
 - ❑ **create role** *teaching_assistant*
 - ❑ **grant** *teaching_assistant* **to instructor**;
 - ▶ *Instructor* inherits all privileges of *teaching_assistant*
- ❑ Chain of roles
 - ❑ **create role** *dean*;
 - ❑ **grant** *instructor* **to dean**;
 - ❑ **grant** *dean* **to Satoshi**;

